

hw02_gap

Avril Metcalfe-Roach

18 September 2019

Exercise 1: Dplyr package

1.1 - Subset gapminder to 3 countries, 1970s.

Countries: Canada, India, Italy

```
filtered <- gapminder %>%  
  filter(country %in% c("Canada", "India", "Italy"),  
         year %in% c(1970:1979))  
knitr::kable(filtered, format="markdown")
```

country	continent	year	lifeExp	pop	gdpPercap
Canada	Americas	1972	72.880	22284500	18970.5709
Canada	Americas	1977	74.210	23796400	22090.8831
India	Asia	1972	50.651	567000000	724.0325
India	Asia	1977	54.208	634000000	813.3373
Italy	Europe	1972	72.190	54365564	12269.2738
Italy	Europe	1977	73.480	56059245	14255.9847

1.2 - Select country, gdpPercap using %>%

```
country_gdp <- filtered %>%  
  select(country, gdpPercap)  
knitr::kable(country_gdp, format="markdown")
```

country	gdpPercap
Canada	18970.5709
Canada	22090.8831
India	724.0325
India	813.3373
Italy	12269.2738
Italy	14255.9847

1.3 - Entries with a negative change in life expectancy from previous line

```
exp_list <- gapminder$lifeExp  
change <- diff(exp_list, lag=1, differences=1)  
# Add NA value to beginning of change vector:  
change_2 <- append(change, NA, after=0)  
# Create new tibble with delta life expectancy as a column:  
gapminder_lifeExp <- gapminder  
gapminder_lifeExp$delta <- change_2  
# Filter for ALL negative changes
```

```
gapminder_redExp <- gapminder_lifeExp %>%
  filter(delta < 0)
```

```
gapminder_redExp
```

```
## # A tibble: 221 x 7
##   country    continent  year lifeExp      pop gdpPercap  delta
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>    <dbl>
## 1 Albania    Europe    1992   71.6  3326498   2497.  -0.419
## 2 Algeria    Africa    1952   43.1  9279525   2449. -33.3
## 3 Angola     Africa    1952   30.0  4232095   3521. -42.3
## 4 Angola     Africa    1987   39.9  7874230   2430. -0.036
## 5 Australia  Oceania    1952   69.1  8691212  10040. -6.20
## 6 Austria    Europe    1952   66.8  6927772   6137. -14.4
## 7 Bahrain    Asia      1952   50.9   120447   9867. -28.9
## 8 Bangladesh Asia      1952   37.5  46886859    684. -38.2
## 9 Benin       Africa    1952   38.2  1738315   1063. -41.2
## 10 Benin      Africa    2002   54.4  7026113   1373. -0.371
## # ... with 211 more rows
```

1.4 - Gapminder: max GDP per capita per country

Original question (using group_by())

```
# Create new column that lists the max GDP per country
```

```
gap_max_gdp <- gapminder %>%
  group_by(country) %>%
  mutate(max_gdp = max(gdpPercap)) %>%
  ungroup()
```

```
# Filters gapminder to only show max GDP; removes the redundant 'max GDP' column
```

```
max_per_country <- gap_max_gdp %>%
  filter(gdpPercap == max_gdp) %>%
  subset(select = -max_gdp)
```

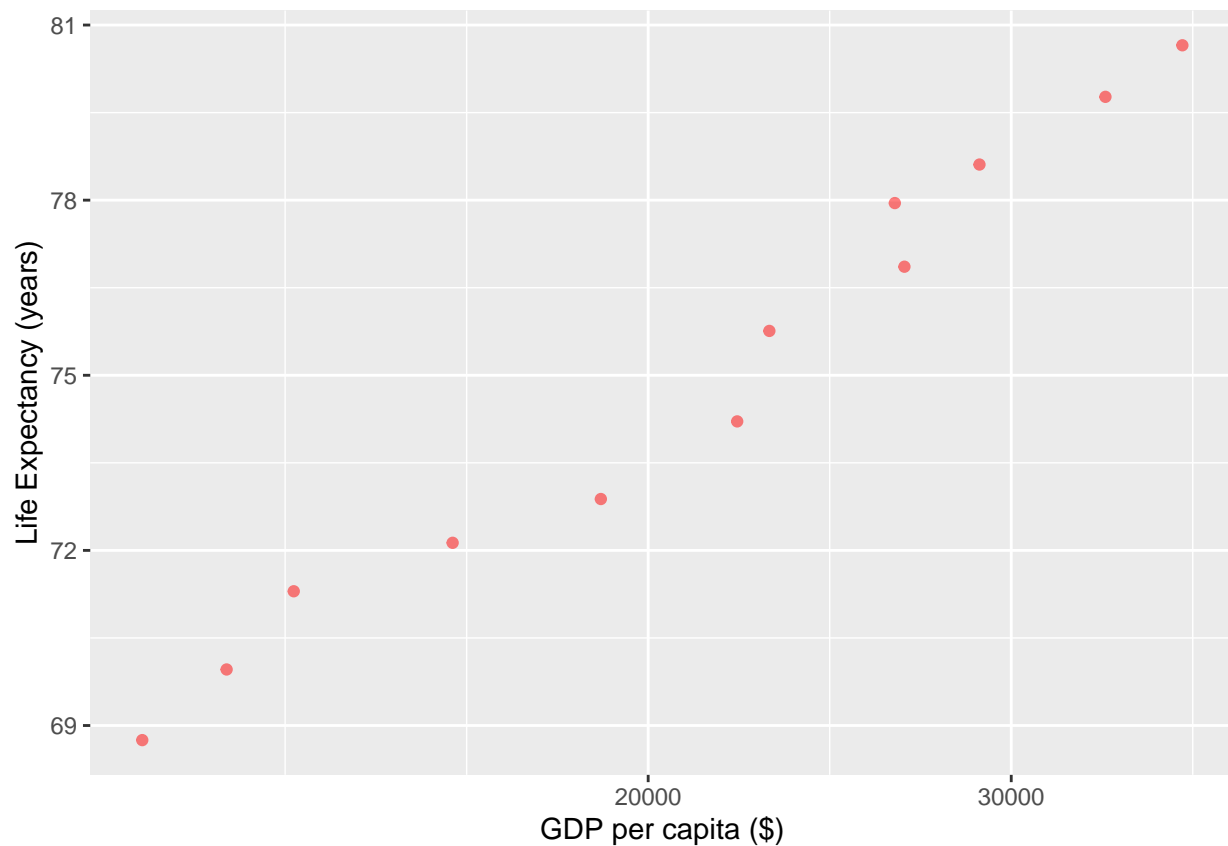
```
max_per_country
```

```
## # A tibble: 142 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1982   39.9  12881816    978.
## 2 Albania     Europe    2007   76.4   3600523   5937.
## 3 Algeria     Africa    2007   72.3  33333216   6223.
## 4 Angola      Africa    1967   36.0   5247469   5523.
## 5 Argentina   Americas  2007   75.3  40301927  12779.
## 6 Australia   Oceania    2007   81.2  20434176  34435.
## 7 Austria     Europe    2007   79.8   8199783   36126.
## 8 Bahrain     Asia      2007   75.6    708573   29796.
## 9 Bangladesh  Asia      2007   64.1  150448339   1391.
## 10 Belgium    Europe    2007   79.4  10392226   33693.
## # ... with 132 more rows
```

1.5 - Canadian Life Expectancy vs GDP

```
# Select data
canadians <- gapminder %>%
  filter(country=="Canada") %>%
  select(lifeExp,gdpPercap)

# Plot in ggplot
ggplot(canadians, aes(gdpPercap,lifeExp)) +
  geom_point(alpha=0.5, colour = "red") +
  scale_x_log10("GDP per capita ($)") +
  ylab("Life Expectancy (years)")
```



Exercise 2: Explore individual variables with dplyr

Categorical variable: continent

Possible range of continent

- Assuming we're not creating any new continents, this variable is inherently limited to the seven continents.
 - Note: North & South America are grouped into 'Americas'
- Possibilities: `c(Asia, Americas, Europe, Africa, Oceania, Antarctica)`
 - Note: Antarctica has no entries in gapminder, as it is a research base.

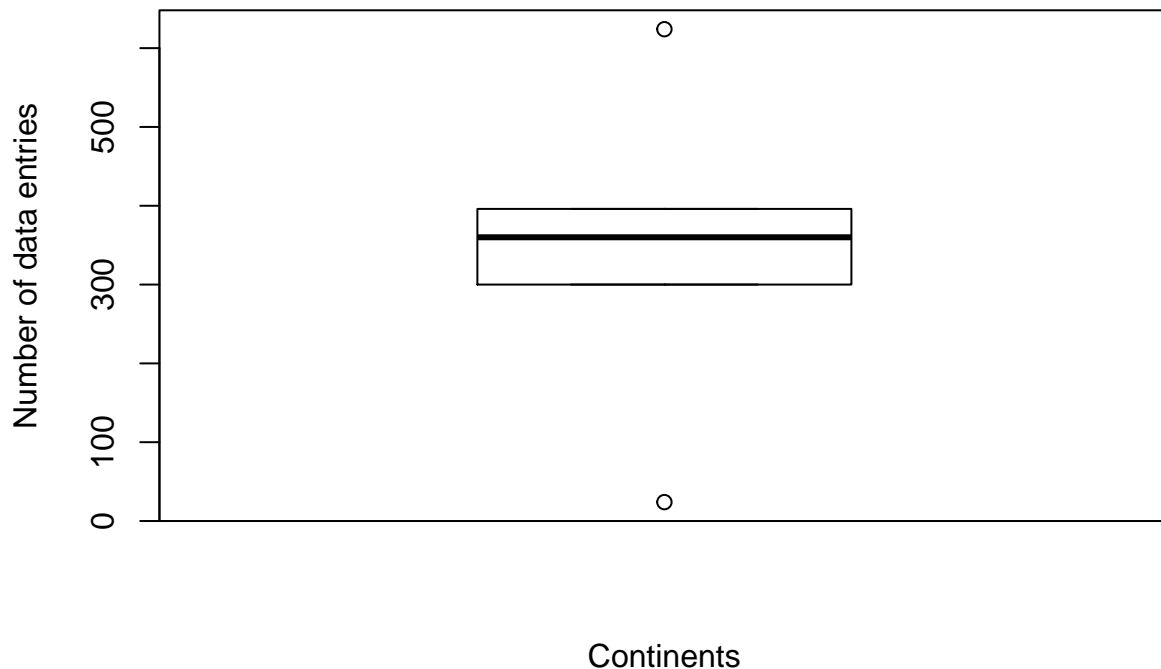
Spread of values

Box and Whisker summary of data:

```
con_only <- gapminder %>%  
  select(continent)  
con_sum <- count(con_only, continent) %>% # dplyr: table of counts per continent  
  as_tibble()  
knitr::kable(summary(con_sum), format="markdown") # Print summary of the data
```

continent	n
Africa :1	Min. : 24.0
Americas:1	1st Qu.:300.0
Asia :1	Median :360.0
Europe :1	Mean :340.8
Oceania :1	3rd Qu.:396.0
NA	Max. :624.0

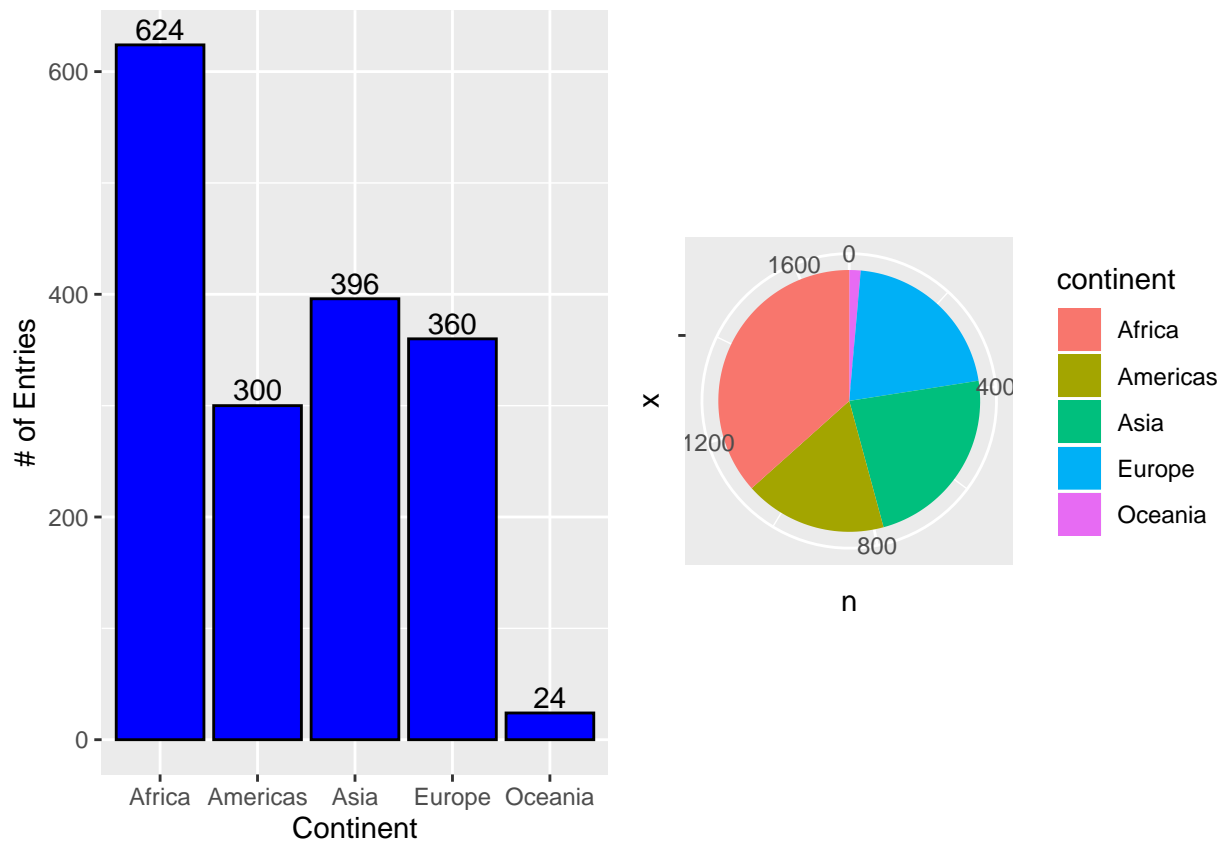
```
#Visualize the summary in a boxplot  
boxplot(con_sum$n,  
  ylab="Number of data entries",  
  xlab = "Continents")
```



The number of datapoints for each populated continent (e.g. Antarctica not included) **ranged from 24 to 624**. The **mean and median were 341 and 360** respectively, with 50% of the data falling between 300 and 396 entries.

Visual representation of distribution:

```
bar_plot <- ggplot(con_sum, aes(continent,n)) +  
  geom_col(colour="black",fill="blue") +  
  geom_text(aes(label=n), vjust=-0.25) +  
  ylab("# of Entries") +  
  xlab("Continent")  
  
pie_plot <- ggplot(con_sum, aes(x='',y=n,fill=continent)) +  
  geom_bar(width=1, stat = "identity") +  
  coord_polar(theta="y")  
  
require(gridExtra) # Arrange them side by side  
  
## Loading required package: gridExtra  
##  
## Attaching package: 'gridExtra'  
## The following object is masked from 'package:dplyr':  
##  
## combine  
gridExtra::grid.arrange(bar_plot,pie_plot,ncol=2)
```



Generally, Oceania is very underrepresented, comprising just 24 out of 1704 entries. Conversely, African data was included at twice the rate of the average at 624 entries. The other three continents are relatively evenly represented.

Quantitative variable: pop (population)

Range of pop

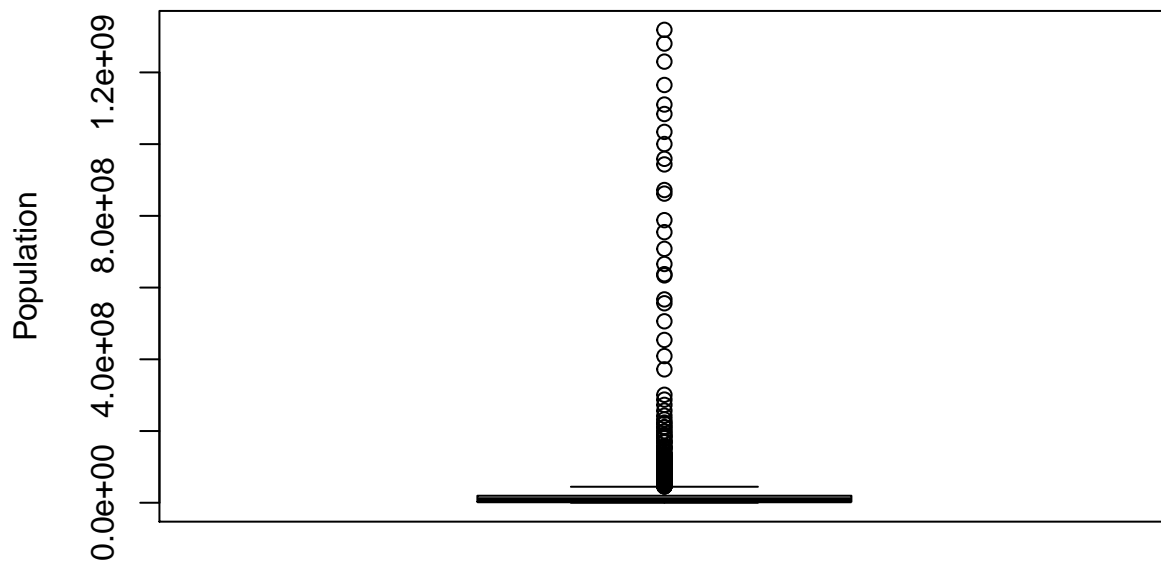
The value of pop must be a Natural (>0) number. No strict upper limit is specified, but should logically be approximately 1.4 billion (the population of China).

Visualizing the spread of all population data:

```
pop_only <- gapminder %>%  
  select(pop)  
knitr::kable(summary(pop_only), format="markdown")
```

pop
Min. :6.001e+04
1st Qu.:2.794e+06
Median :7.024e+06
Mean :2.960e+07
3rd Qu.:1.959e+07
Max. :1.319e+09

```
boxplot(pop_only$pop,  
        ylab="Population",  
        xlab = "")
```



As demonstrated by the boxplot, the vast majority of the data (all data within the whiskers/confidence interval) comprise a tiny fraction of the possible range of population values. **50% of the data describes a**

population between 2.8-19.6 million, with the median population being 7 million. The average is much higher at 29.6 million as the high-population outliers are skewing the data. The minimum and maximum populations are 60 000 and 1.32 billion respectively.

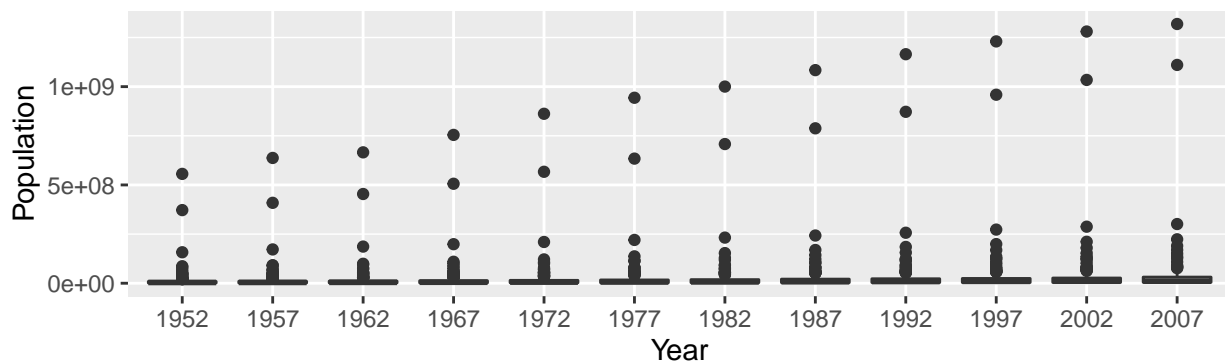
There are 12 entries for each country, as they were sampled at every time point. We can divide the data by year to see how the average populations change over time:

```
# Linear Plot
pop_time <- gapminder %>% # gapminder is gapminder in tibble format
  select(year, pop)
pop_time$year <- as.factor(pop_time$year) # Only factors can be used to plot side-by-side boxplots
pop_time_plot <- ggplot(pop_time, aes(year, pop)) +
  geom_boxplot() +
  xlab("Year") +
  ylab("Population") +
  ggtitle("Linear")

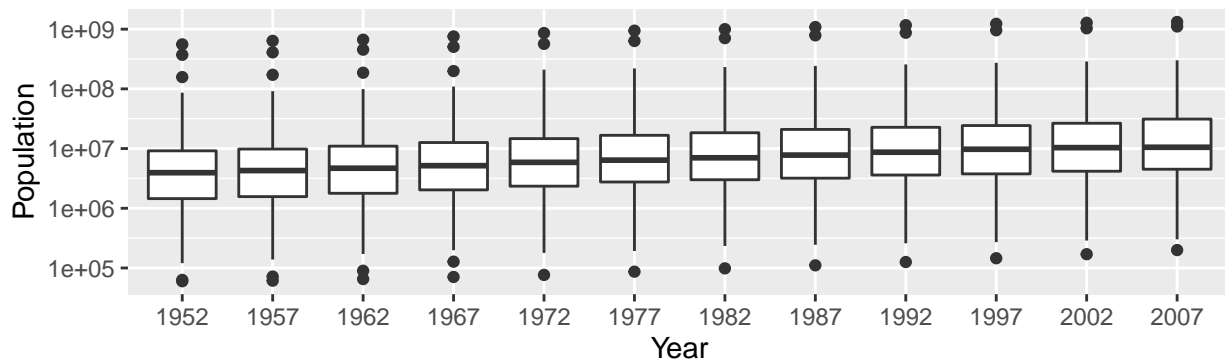
# Log Transformed
log_plot <- pop_time_plot +
  scale_y_log10() +
  ggtitle("Logarithmic")

# Side by Side Output
require(gridExtra)
gridExtra::grid.arrange(pop_time_plot, log_plot, nrow=2)
```

Linear



Logarithmic



The above graph makes it easier to see that there are only a couple of countries that have populations significantly outside of the statistical range. In the linear plot, the significant population size and fast

growth of China and India in particular make the population growth of the rest of the world less apparent. By transforming the y axis to a log 10 scale, all of the outliers can be captured and the general trends become apparent: the **IQR (middle 50%) of the data moves up the y axis, showing exponential population growth.**

Exercise 3: Plot Exploration

Scatterplot of [CO2]ambient vs [CO2]uptake

Dataset: *CO2 - Carbon Dioxide Uptake in Grass Plants*

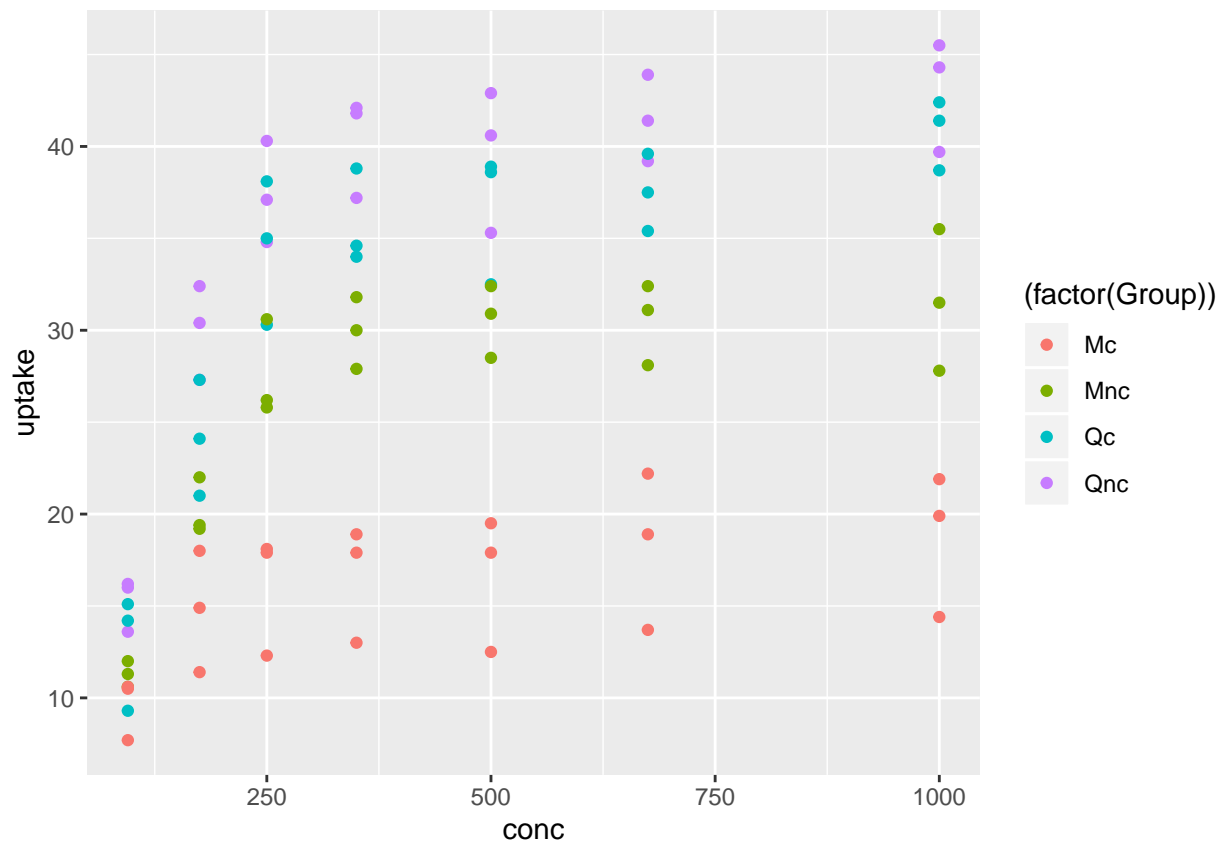
Here are the main parameters of the CO2 dataset:

```
knitr::kable(summary(CO2),format="markdown")
```

Plant	Type	Treatment	conc	uptake
Qn1 : 7	Quebec :42	nonchilled:42	Min. : 95	Min. : 7.70
Qn2 : 7	Mississippi:42	chilled :42	1st Qu.: 175	1st Qu.:17.90
Qn3 : 7	NA	NA	Median : 350	Median :28.30
Qc1 : 7	NA	NA	Mean : 435	Mean :27.21
Qc3 : 7	NA	NA	3rd Qu.: 675	3rd Qu.:37.12
Qc2 : 7	NA	NA	Max. :1000	Max. :45.50
(Other):42	NA	NA	NA	NA

Plant is the type of plant, Type is the location of the plant, conc is the ambient CO2 concentration, and uptake is the CO2 absorbed by the plant.

```
CO2 %>%
  mutate(Group = (if_else
    (Type=="Quebec",
      if_else
        (Treatment=="nonchilled","Qnc","Qc"),
        if_else(Treatment=="nonchilled","Mnc","Mc")))) %>%
  ggplot(aes(conc,uptake)) +
  geom_point(aes(colour=(factor(Group))))
```

The above graph shows the change in CO2 uptake as a function of CO2 ambient concentration. In the legend, M/Q denote location (Mississippi vs Quebec) and c/nc denote treatment (chilled/not chilled).

From the plot, it can be seen that Quebec plants are more efficient at carbon fixation than Mississippi plants, and chilling the plants reduces their efficiency.

Graph #2

Dataset: **esoph** - *Smoking, Alcohol and (O)esophageal Cancer*

Here are the main parameters of the CO2 dataset:

```
esoph_cancer <- esoph %>%
  rename("Alcohol Intake"=alcgp) %>%
  rename("Tobacco Intake"=tobgp)

knitr::kable(summary(esoph_cancer),format="markdown")
```

agegp	Alcohol Intake	Tobacco Intake	ncases	ncontrols
25-34:15	0-39g/day:23	0-9g/day:24	Min. : 0.000	Min. : 1.00
35-44:15	40-79 :23	10-19 :24	1st Qu.: 0.000	1st Qu.: 3.00
45-54:16	80-119 :21	20-29 :20	Median : 1.000	Median : 6.00
55-64:16	120+ :21	30+ :20	Mean : 2.273	Mean :11.08
65-74:15	NA	NA	3rd Qu.: 4.000	3rd Qu.:14.00
75+ :11	NA	NA	Max. :17.000	Max. :60.00

Note: Ncases and ncontrols summaries are meaningless, as they supply weightings to the alcohol and tobacco groups.

We will be looking at the **rate of esophageal cancer relative to alcohol consumption** and irrespective of age.

```
# Collapse data: sum cases/controls by alcohol group
p2 <- aggregate(cbind(esoph_cancer$ncases,esoph_cancer$ncontrols),
               by=list(esoph_cancer$`Alcohol Intake`),
               FUN=sum)

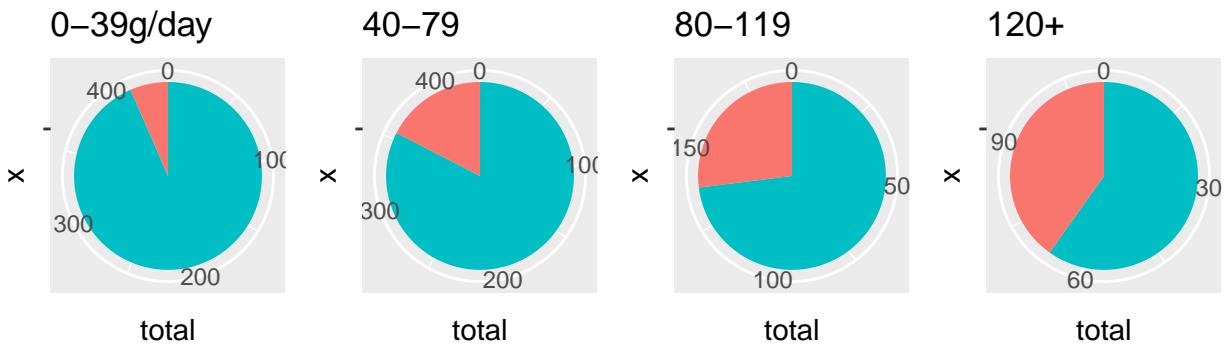
cases <- mutate(p2,status="case") %>%
  rename("total"=V1) %>%
  select(-V2)
ctrls <- mutate(p2,status="ctrl") %>%
  rename("total"=V2) %>%
  select(-V1)
all_data <- merge(cases,ctrls,all=TRUE)

all_39 <- filter(all_data,Group.1=="0-39g/day")
all_79 <- filter(all_data,Group.1=="40-79")
all_119 <- filter(all_data,Group.1=="80-119")
all_120 <- filter(all_data,Group.1=="120+")

esoph_pie <- function(df){
  ggplot(df, aes(x='',y=total,fill=status,)) +
    geom_bar(width=1,stat="identity") +
    coord_polar(theta="y") +
    guides(fill = FALSE, color = FALSE, linetype = FALSE, shape = FALSE) +
    ggtitle(df$Group.1[1])
}

g1 <- esoph_pie(all_39)
g2 <- esoph_pie(all_79)
g3 <- esoph_pie(all_119)
g4 <- esoph_pie(all_120)

require(gridExtra)
gridExtra::grid.arrange(g1,g2,g3,g4,ncol=4)
```



The esophageal cancer cases are represented by pink, while the controls are in blue. Titles indicate the number of grams of alcohol ingested per day. As the alcohol intake increases, so do the overall cancer rate.

Recycling

Code in question:

```
filter(gapminder, country == c("Rwanda", "Afghanistan"))
```

```
## # A tibble: 12 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int> <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1957  30.3  9240934  821.
## 2 Afghanistan Asia      1967  34.0 11537966  836.
## 3 Afghanistan Asia      1977  38.4 14880372  786.
## 4 Afghanistan Asia      1987  40.8 13867957  852.
## 5 Afghanistan Asia      1997  41.8 22227415  635.
## 6 Afghanistan Asia      2007  43.8 31889923  975.
## 7 Rwanda     Africa    1952   40   2534927  493.
## 8 Rwanda     Africa    1962   43   3051242  597.
## 9 Rwanda     Africa    1972  44.6  3992121  591.
## 10 Rwanda     Africa    1982  46.2  5507565  882.
## 11 Rwanda     Africa    1992  23.6  7290203  737.
## 12 Rwanda     Africa    2002  43.4  7852401  786.
```

Though at first glance this code appears to select the data for Afghanistan and Rwanda, but **half of the data is missing**. Every other entry is omitted, in a way where both countries are never represented in the

same year. (ex. Rwanda 1952, Afghanistan 1957...)

Because of the inclusion of a list, the function will go along the year-sorted data, taking turns selecting each country for each value of year. The code can be fixed by using an 'or' statement:

```
filter(gapminder, country == "Rwanda" | country == "Afghanistan")
```

```
## # A tibble: 24 x 6
##   country    continent  year lifeExp      pop gdpPercap
##   <fct>      <fct>    <int>  <dbl>    <int>    <dbl>
## 1 Afghanistan Asia      1952   28.8  8425333    779.
## 2 Afghanistan Asia      1957   30.3  9240934    821.
## 3 Afghanistan Asia      1962   32.0 10267083    853.
## 4 Afghanistan Asia      1967   34.0 11537966    836.
## 5 Afghanistan Asia      1972   36.1 13079460    740.
## 6 Afghanistan Asia      1977   38.4 14880372    786.
## 7 Afghanistan Asia      1982   39.9 12881816    978.
## 8 Afghanistan Asia      1987   40.8 13867957    852.
## 9 Afghanistan Asia      1992   41.7 16317921    649.
## 10 Afghanistan Asia      1997   41.8 22227415    635.
## # ... with 14 more rows
```

In this way, all of the data are collected.

Tibbles

All small tables as `knitr::kable()`. `DT::datatable()` only works in html; therefore, the .Rmd and html file on github contains `datatable()` format for large tables, but `as_tibble()` is used in the .md file.